

Indexing HTML files in Solr¹

This tutorial explains how to index html files in Solr using the built-in post tool, which leverages Apache Tika and auto extracts content from html files. You should have already downloaded and installed Solr, see

<http://www.scf.usc.edu/~csci572/2016Spring/hw3/SolrInstallation.pdf>

Apache Tika is a library that is used for document type detection and content extraction from various file formats. Internally, Tika uses various existing document parsers and document type detection techniques to detect and extract data. For example, in case of HTML pages, Tika uses its HTMLParser to strip out all html tags and only stores the content from the html pages. Tika is a powerful tool and is very useful when you have crawled various types of documents, e.g. PDF, Images, Videos. Tika is included with the Solr installation.

Using post tool and TIKa

1. Start the Solr server, cd into solr-5.3.1 folder and enter the command: **bin/solr start**
2. Create a new core, in this tutorial the core name is **myexample** using the command: **bin/solr create -c myexample**
After the core is created you should be able to see the following message. It also displays where your instance directory is placed. In this case it is created in solr-5.3.1/server/solr/ with the name myexample.

```
Setup new core instance directory:
/home/solr-5.3.1/server/solr/myexample

Creating new core 'myexample' using command:
http://localhost:8983/solr/admin/cores?action=CREATE&name=myexample&instanceDir=myexample

{
  "responseHeader":{
    "status":0,
    "QTime":1213},
  "core":"myexample"}
```

3. Let us examine the directory structure of this folder. cd into the folder myexample in server/solr. There will be a **conf** folder, **data** folder and a **core.properties** file. The conf folder contains the schema.xml file and the solrconfig.xml file, these are the files that have to be modified for specific field level analysis during indexing and querying. For example you can specify that a field from your document has to be tokenized using white space delimiter, whereas another field has to be tokenized into n-grams, depending on how you want to index and later query these fields. You can specify stop words in stopwords.txt file which would enable Solr to eliminate these words during indexing. The data folder contains the indexed data and logs generated by Solr. The core.properties file contains metadata information about the core, e.g, the core name, the directory name where the core data is stored.

conf core.properties data

4. Solr inherently uses Tika for extracting content from the documents that will be indexed. Tika uses the TagSoup library to support virtually any kind of HTML found on the web. The output from the HtmlParser class is used as the streamed content for indexing into Solr. Before indexing html files, we have to edit the schema file to make sure that all the text content from the html pages extracted by Tika are mapped correctly. To do this, go to the conf folder of the core "myexample". You will see a file named – "**managed-schema**". First rename the file to **schema.xml**. Now open the xml to edit it. Below is a portion of the schema.xml file.

¹ This tutorial was authored by Majisha Parambath (with a little help from Prof. Horowitz)

```

<field name="id" type="string" indexed="true" stored="true" required="true" multiValued="false" />
<field name="_version_" type="long" indexed="true" stored="true"/>
<field name="_root_" type="string" indexed="true" stored="false"/>
<field name="_text_" type="text_general" indexed="true" stored="false" multiValued="true"/>
<copyField source="*" dest="_text_" />

```

Fields are defined in the field element of schema.xml as shown in the above figure. These are field definitions for the documents that you will be indexing, and each document that gets indexed will have these fields. You can define various properties for the fields, such as “indexed” if set to true, then the content in this field will be indexed, “stored” if set then the content of this field is stored in Solr and we can retrieve this from the query responses, “required” indicates whether this field is mandatory in a document for indexing, multiValued is true meaning it can appear multiple times in a document. In the above example, an “id” field, “_version_” and “_root_” fields are auto generated while indexing. These may not necessarily be part of your original documents.

Solr has a mechanism for making copies of fields so that you can apply several distinct field types to a single piece of incoming information. You can see that the copyField element copies all the fields into the destination field “_text_”. The “_text_” field is also only indexed and not stored because there is no need to have redundant data. Basically what we are trying to accomplish here, is to combine all the information from different fields into a single field, “_text_”.

5. There is a commented section in the schema.xml file, which has the comment “Text fields from the SolrCell¹ to search by default in our catch-all field” as shown in the (already edited) figure below.

```

<!-- Text fields from SolrCell to search by default in our catch-all field -->
<copyField source="title" dest="_text_" />
<copyField source="author" dest="_text_" />
<copyField source="description" dest="_text_" />
<copyField source="keywords" dest="_text_" />
<copyField source="content" dest="_text_" />
<copyField source="content_type" dest="_text_" />
<copyField source="resourcename" dest="_text_" />
<copyField source="url" dest="_text_" />

```

These are basically some of the fields extracted by TIKa using the HTML Parser class. TIKa automatically identifies the title, author, description, keywords from the metadata of the documents. The content field consists of all the text that the TIKa Parser could extract from rich documents such as Word, PDF etc. In case of HTML files, TIKa’s HTML Parser extracts all the text available on the HTML page and stores them in the content field. Uncomment this section, and change the destination to the previously defined “_text_” field as shown in the figure above.

6. Save your changes and go back to the home directory of Solr. Let us try to index some html pages. In this example, I have crawled the webpages from losangeles.eventful.com and stored the crawled pages in a folder named “crawl_data”. You should be using the folder where you have crawled all the webpages as part of your crawling assignment. Indexing is performed by using the command:

```
bin/post -c <core_name> -filetypes html <path_to_crawl_folder>/
```

The filetypes option specified above implies we are only indexing html files. If we are indexing various document types together, we can ignore the filetypes option, then the command will be simply:

```
bin/post -c <core_name> <path_to_crawl_folder>/
```

The crawl folder in this example is “crawl_data”, the core to which I am indexing is “myexample” and I am going to index only html files, therefore the command will be:

```
bin/post -c myexample -filetypes html crawl_data/
```

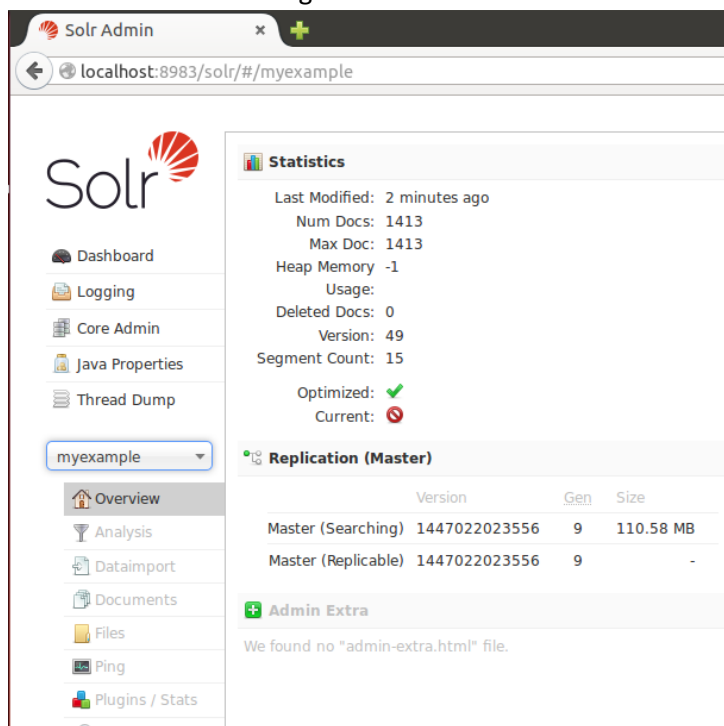
You should see output similar to the one below if the command was successful:

```
SimplePostTool version 5.0.0
Posting files to [base] url http://localhost:8983/solr/myexample/update...
Entering auto mode. File endings considered are html
Entering recursive mode, max depth=999, delay=0s
Indexing directory crawl_data (1413 files, depth=0)
POSTing file rolling-robots-1.html (text/html) to [base]/extract
POSTing file gl-tatuaje.html (text/html) to [base]/extract
POSTing file fitness-joq.html (text/html) to [base]/extract
POSTing file cafe-del-sol-3.html (text/html) to [base]/extract
POSTing file teaglad-8-los-angeles.html (text/html) to [base]/extract
POSTing file sumbody-studio-studio-city.html (text/html) to [base]/extract
POSTing file pilates-plus-wv-los-angeles.html (text/html) to [base]/extract
POSTing file aziam-yoga-and-surf-hikes.html (text/html) to [base]/extract
POSTing file crossfit-anaheim-2.html (text/html) to [base]/extract
POSTing file harvelle-s-1.html (text/html) to [base]/extract
POSTing file market-place-restaurant.html (text/html) to [base]/extract
POSTing file yanks-air-museum-los-angeles.html (text/html) to [base]/extract
```

After all the files are indexed, solr will auto commit the index and you should see the following output.

```
1413 files indexed.
COMMITting Solr index changes to http://localhost:8983/solr/myexample/update...
Time spent: 0:01:52.729
```

7. So now we can check the Solr UI and see if the files have gotten indexed. Open the browser and go to <http://localhost:8983/solr/>. Select the core “myexample” from the dropdown. You should see an output showing the statistics similar to the figure below.



The screenshot shows the Solr Admin interface in a web browser. The address bar shows `localhost:8983/solr/#/myexample`. The left sidebar contains navigation links: Dashboard, Logging, Core Admin, Java Properties, Thread Dump, and a dropdown menu currently showing 'myexample'. Below the dropdown are links for Overview, Analysis, Dataimport, Documents, Files, Ping, and Plugins / Stats. The main content area displays the 'Statistics' section for the 'myexample' core. It shows the following data:

- Last Modified: 2 minutes ago
- Num Docs: 1413
- Max Doc: 1413
- Heap Memory: -1
- Usage:
- Deleted Docs: 0
- Version: 49
- Segment Count: 15
- Optimized: ✔
- Current: ✘

Below the statistics is the 'Replication (Master)' section, which shows a table with the following data:

	Version	Gen	Size
Master (Searching)	1447022023556	9	110.58 MB
Master (Replicable)	1447022023556	9	-

At the bottom of the main content area is the 'Admin Extra' section, which states: 'We found no "admin-extra.html" file.'

6. Let us see how the TIKA parser parsed the html pages and what fields were created by using TIKA. Select the query option, and just submit the default query (*:*). Below is a snapshot of a portion of one of the web pages from `losangeles.events.com` followed by the results produced by the TIKA parser. Some notable fields extracted here are title, latitude, longitude, description etc. We can see the HTML page source from where the fields were extracted below.

```

<meta name="description" content="LA Gift & Home Market on Jan 28, 2016 in Los Angeles, CA at California Market Center." />

<link rel="canonical" href="http://losangeles.eventful.com/events/la-gift-home-market-/E0-001-086391282-0" />
<link rel="bookmark" href="http://losangeles.eventful.com/events/la-gift-home-market-/E0-001-086391282-0" />

<!-- meta: Open Graph -->
<meta property="og:title" content="LA Gift & Home Market" />
<meta property="og:type" content="eventful:event" />
<meta property="og:url" content="http://losangeles.eventful.com/events/la-gift-home-market-/E0-001-086391282-0" />
<meta property="og:image" content="http://s1.evcdn.com/images/block200/fallback/event/categories/conference/conference_default_3.jpg" />

<meta property="og:description" content="LA Gift & Home Market at California Market Center on Thursday Jan 28, 2016 at 12:00AM" />

```

The following figure is snipped from the response of the default query in the SolrUI. Notice that the response uses the JSON format. Also notice that all meta properties in the html are preserved. Notice that some key : value pairs are autogenerated by Tika, e.g. content_encoding, content_type, x_parsed_by, stream_size, and stream_content_type.

```

{
  "id": "/home/solr-5.3.1/CrawlData/E0-001-086391282-0.html",
  "og_type": ["eventful:event"],
  "viewport": ["width=1010, user-scalable=yes"],
  "og_title": ["LA Gift & Home Market"],
  "application_name": ["Eventful"],
  "og_description": ["LA Gift & Home Market at California Market Center on Thursday Jan 28, 2016 at 12:00AM"],
  "geo_lat": [34.0522],
  "icbm": ["34.0522, -118.243"],
  "x_parsed_by": ["org.apache.tika.parser.DefaultParser",
    "org.apache.tika.parser.html.HtmlParser"],
  "fb_app_id": [294833066685],
  "og_site_name": ["Eventful"],
  "og_image": ["http://s1.evcdn.com/images/block200/fallback/event/categories/conference/conference_default_3.jpg"],
  "og_url": ["http://losangeles.eventful.com/events/la-gift-home-market-/E0-001-086391282-0"],
  "msapplication_tileimage": ["http://s1.evcdn.com/store/skin/header/96f5100b-08f8-4652-b2d6-7a88449ca94a.png"],
  "msapplication_tilecolor": ["#000000"],
  "title": ["LA Gift & Home Market in Los Angeles, CA - Jan 28, 2016 12:00 AM | Eventful"],
  "stream_content_type": ["text/html"],
  "description": ["LA Gift & Home Market on Jan 28, 2016 in Los Angeles, CA at California Market Center."],
  "geo_long": [-118.243],
  "stream_size": [81919],
  "content_encoding": ["UTF-8"],
  "dc_title": ["Eventful: LA Gift & Home Market",
    "LA Gift & Home Market in Los Angeles, CA - Jan 28, 2016 12:00 AM | Eventful"],
  "content_type": ["text/html; charset=UTF-8"],
  "resourcename": ["/home/solr-5.3.1/CrawlData/E0-001-086391282-0.html"],
  "_version_": 151731435645504768}

```

TIKA extracted some fields automatically from these pages.

7. We can also query using the hidden field “_text_”. In the following example I am going to search in the _text_ field for “French”.

Request-Handler (qt)

/select

— common —

q

text:french

The query returned the following results. Looking at the JSON we see 55 documents were matched and returned.

```
"response": {"numFound": 55, "start": 0, "docs": [
  {
    "id": "/home/solr-5.3.1/crawl_data/french-market-cafe.html",
    "twitter_image_height": [300],
    "og_latitude": [34.0522],
    "og_type": ["groupon:deal"],
    "keywords": ["Los Angeles, undefined, local deals, all deals, daily deals, coupons, discounts, things to do, groupon"],
    "og_title": ["Cafe Food and Drinks at French Market Cafe (Up to 40% Off). Two Options Available. "],
    "twitter_title": ["Cafe Food and Drinks at French Market Cafe (Up to 40% Off). Two Options Available. "],
    "og_description": ["French-style cafe that draws many European expats serves up Le Cheval ham, swiss, and fried-egg sandwich"],
    "twitter_image_width": [440],
    "og_site_name": ["Groupon"],
    "x_parsed_by": ["org.apache.tika.parser.DefaultParser",
      "org.apache.tika.parser.html.HtmlParser"],
    "fb_app_id": [7829106395],
    "csrf_token": ["1GmWRUDJ-iynkf7HF7ZRQ3xJ8rGEhXpuBATs"],
    "og_image": ["https://img.grouponcdn.com/deal/ktjo4LzUuXbkMMCX39dw/PS-700x420/v1/t440x300.jpg"],
    "og_url": ["https://www.groupon.com/deals/french-market-cafe"],
    "og_longitude": [-118.243],
    "title": ["Cafe Food and Drinks - French Market Cafe | Groupon"],
    "stream_content_type": ["text/html"],
    "stream_size": [123429],
    "description": ["Cafe Food and Drinks at French Market Cafe (Up to 40% Off). Two Options Available. "],
    "content_encoding": ["UTF-8"],
    "twitter_description": ["French-style cafe that draws many European expats serves up Le Cheval ham, swiss, and fried-egg sandwich"],
    "twitter_site": ["groupon"],
    "twitter_image_src": ["https://img.grouponcdn.com/deal/ktjo4LzUuXbkMMCX39dw/PS-700x420/v1/t440x300.jpg"],
    "content_type": ["text/html; charset=UTF-8"],
    "resourcename": ["/home/solr-5.3.1/crawl_data/french-market-cafe.html"],
    "dc_title": ["Cafe Food and Drinks - French Market Cafe | Groupon"],
    "_version_": 1517974795196563456,
  }
]
```

- We can configure Solr to default querying from the field `_text_` by defining the default field in the requestHandler in the `solrconfig.xml` file.

```
<requestHandler name="/select" class="solr.SearchHandler">
  <!-- default values for query parameters can be specified, these
       will be overridden by parameters in the request
  -->
  <lst name="defaults">
    <str name="echoParams">explicit</str>|
    <int name="rows">10</int>
    <!-- <str name="df">text</str> -->
  </lst>
```

Uncomment the str element with name “df”, and replace it in the following manner, where we specify that the default query field has to be “_text_”.

```
<requestHandler name="/select" class="solr.SearchHandler">
  <!-- default values for query parameters can be specified, these
       will be overridden by parameters in the request
  -->
  <lst name="defaults">
    <str name="echoParams">explicit</str>
    <int name="rows">10</int>
    <str name="df">_text_</str>
  </lst>
```

You can reload your core and query with the new configuration.

Note: in practice we would want to build a more appropriate user interface, one in which the query results are properly displayed. Our new user interface would make AJAX calls on the Solr index rather than use the SolrUI.