HW3 Overview

There are 4 components to this homework; you will possibly not need all of them;

1. Installing Ubuntu
2. Installing Solr
3. Using Solr to Index your downloaded web pages
4. the exercise - comparing ranking algorithms

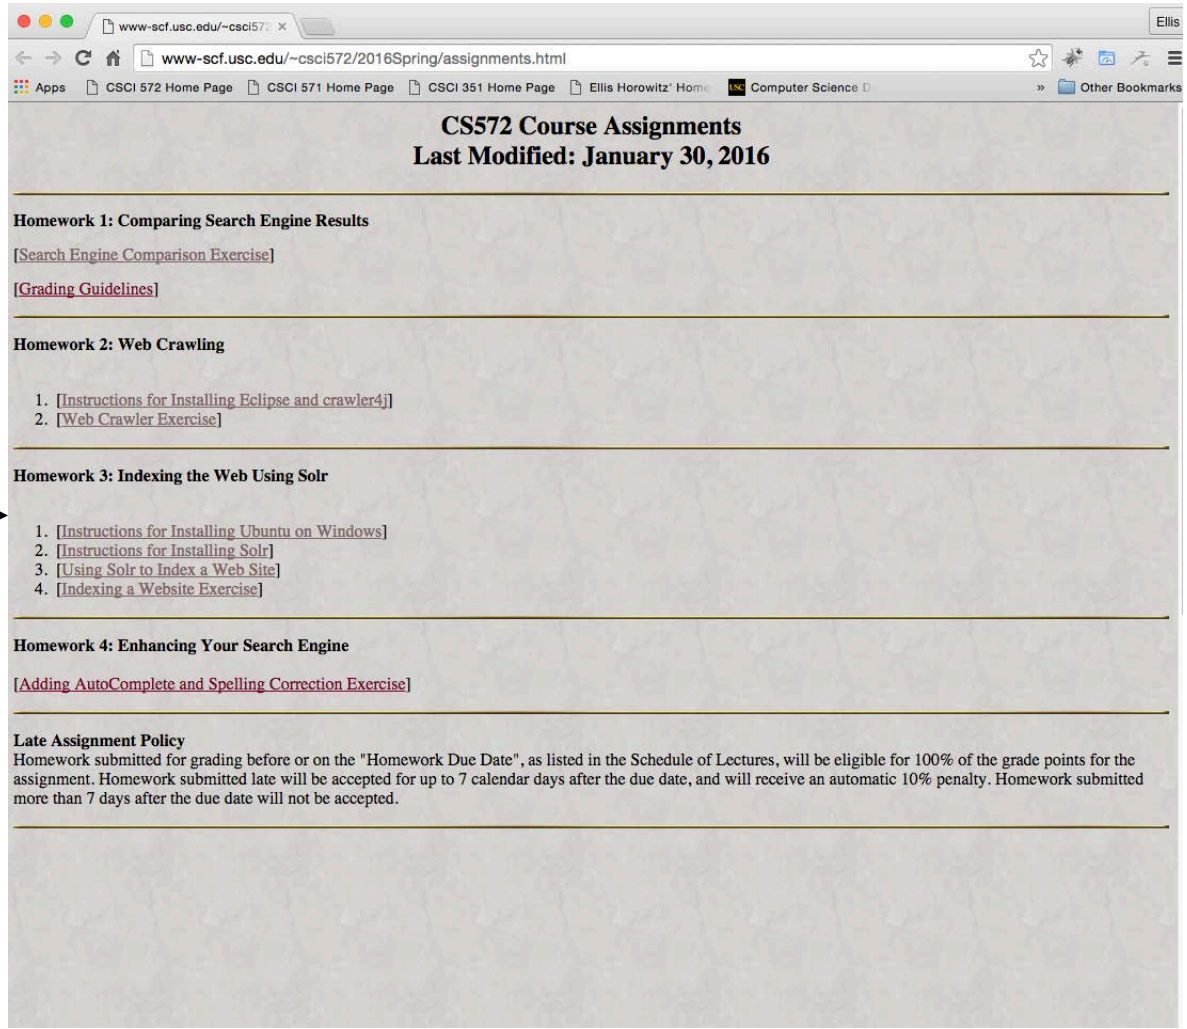

Step 1: Ubuntu with VirtualBox

- **VirtualBox** is an open source, freely available Windows application (it also runs on other platforms) that lets you run multiple operating systems on your single machine
  - E.g. run Windows on a Mac, run Linux on Windows
  - Major supported operating systems include: Windows NT 4.0, Windows 2000, Windows 8, Windows 10, DOS Windows 3.x, Linux, Solaris, FreeBSD, OpenBSD
- **Ubuntu** is a Linux-based operating system distributed on personal computers, smartphones and network servers. It uses **Unity** as its default desktop environment
- **Solr requires a Unix environment** to run, so step 1 is required if you plan to use your Windows laptop

Step 1: Setting Up Ubuntu with VirtualBox

1. Download the free version of VirtualBox for Windows machines
   - Instructions can be found here

   http://www-scf.usc.edu/~csci572/Exercises/UbuntuVirtualBoxFinal.pdf
2. Download the Ubuntu 64-bit version
3. Run VirtualBox and select your Ubuntu version as the New Application
4. Set various parameters
5. Install Ubuntu and you should be ready to run

Your Ubuntu/Unity Desktop

Built-in applications including Firefox browser

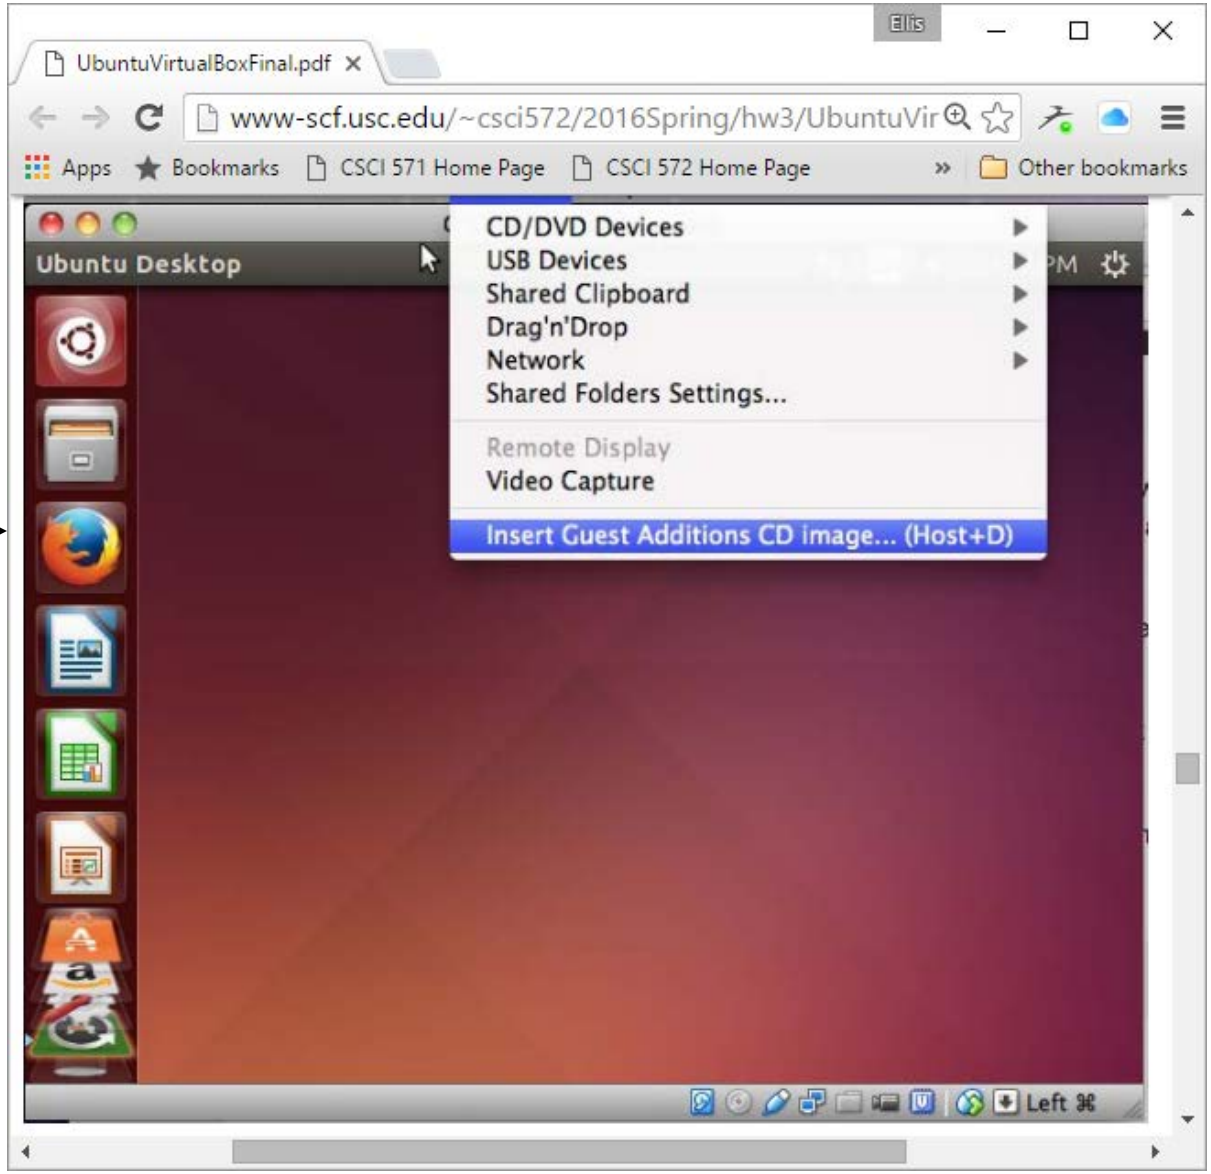

Step 2: Installing Solr

- Solr is an open source enterprise search server based on the Lucene Java search library
- Instructions for downloading and installing Solr can be found here
  - http://www-scf.usc.edu/~csci572/2016Spring/hw3/SolrInstallation.pdf
- Fast, high performance, scalable search/Information Retrieval library
- Initially developed by Doug Cutting (Also author of Hadoop)
- it provides for Indexing and Searching of documents
- produces an Inverted Index of documents
- Provides advanced Search options like synonyms, stopwords, based on similarity, proximity.
- **http://lucene.apache.org/ is the main page for both Lucene and Solr**

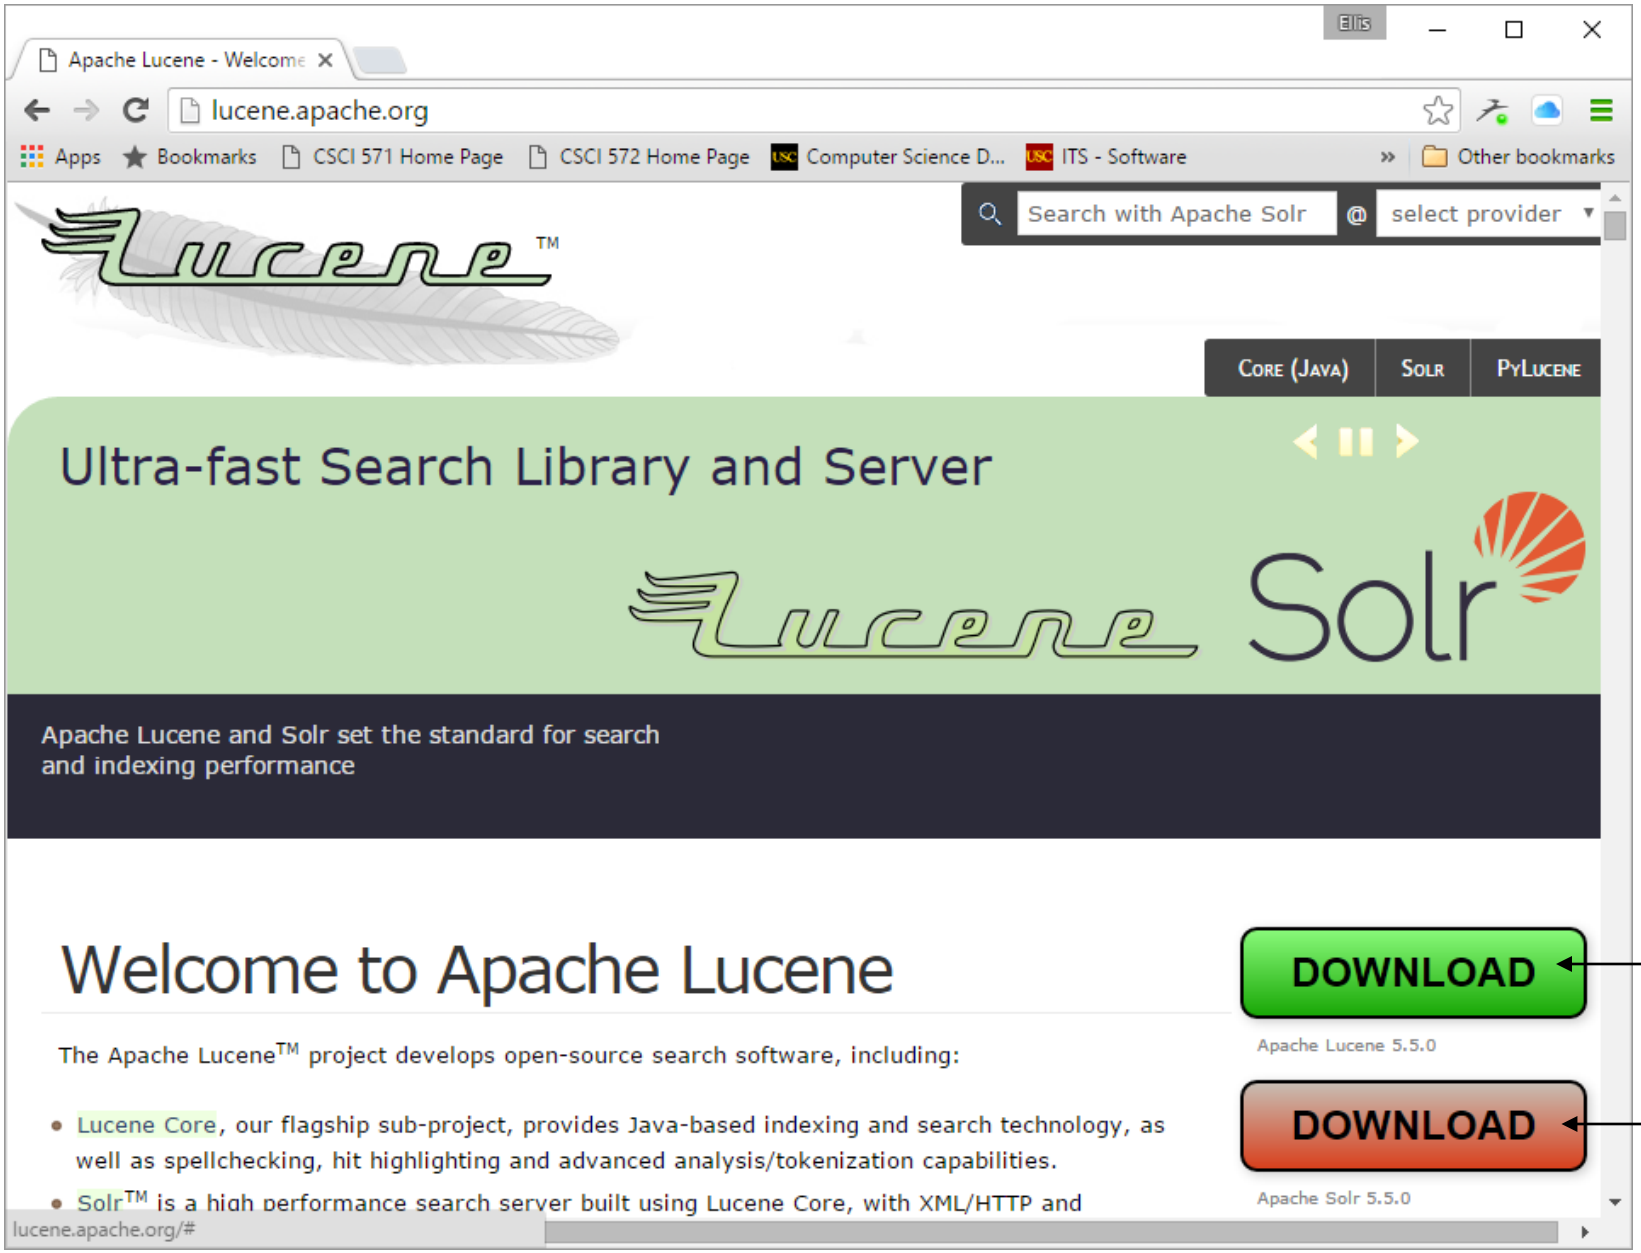

Download Lucene

Download Solr

Lucene Internals - Positional Inverted Index

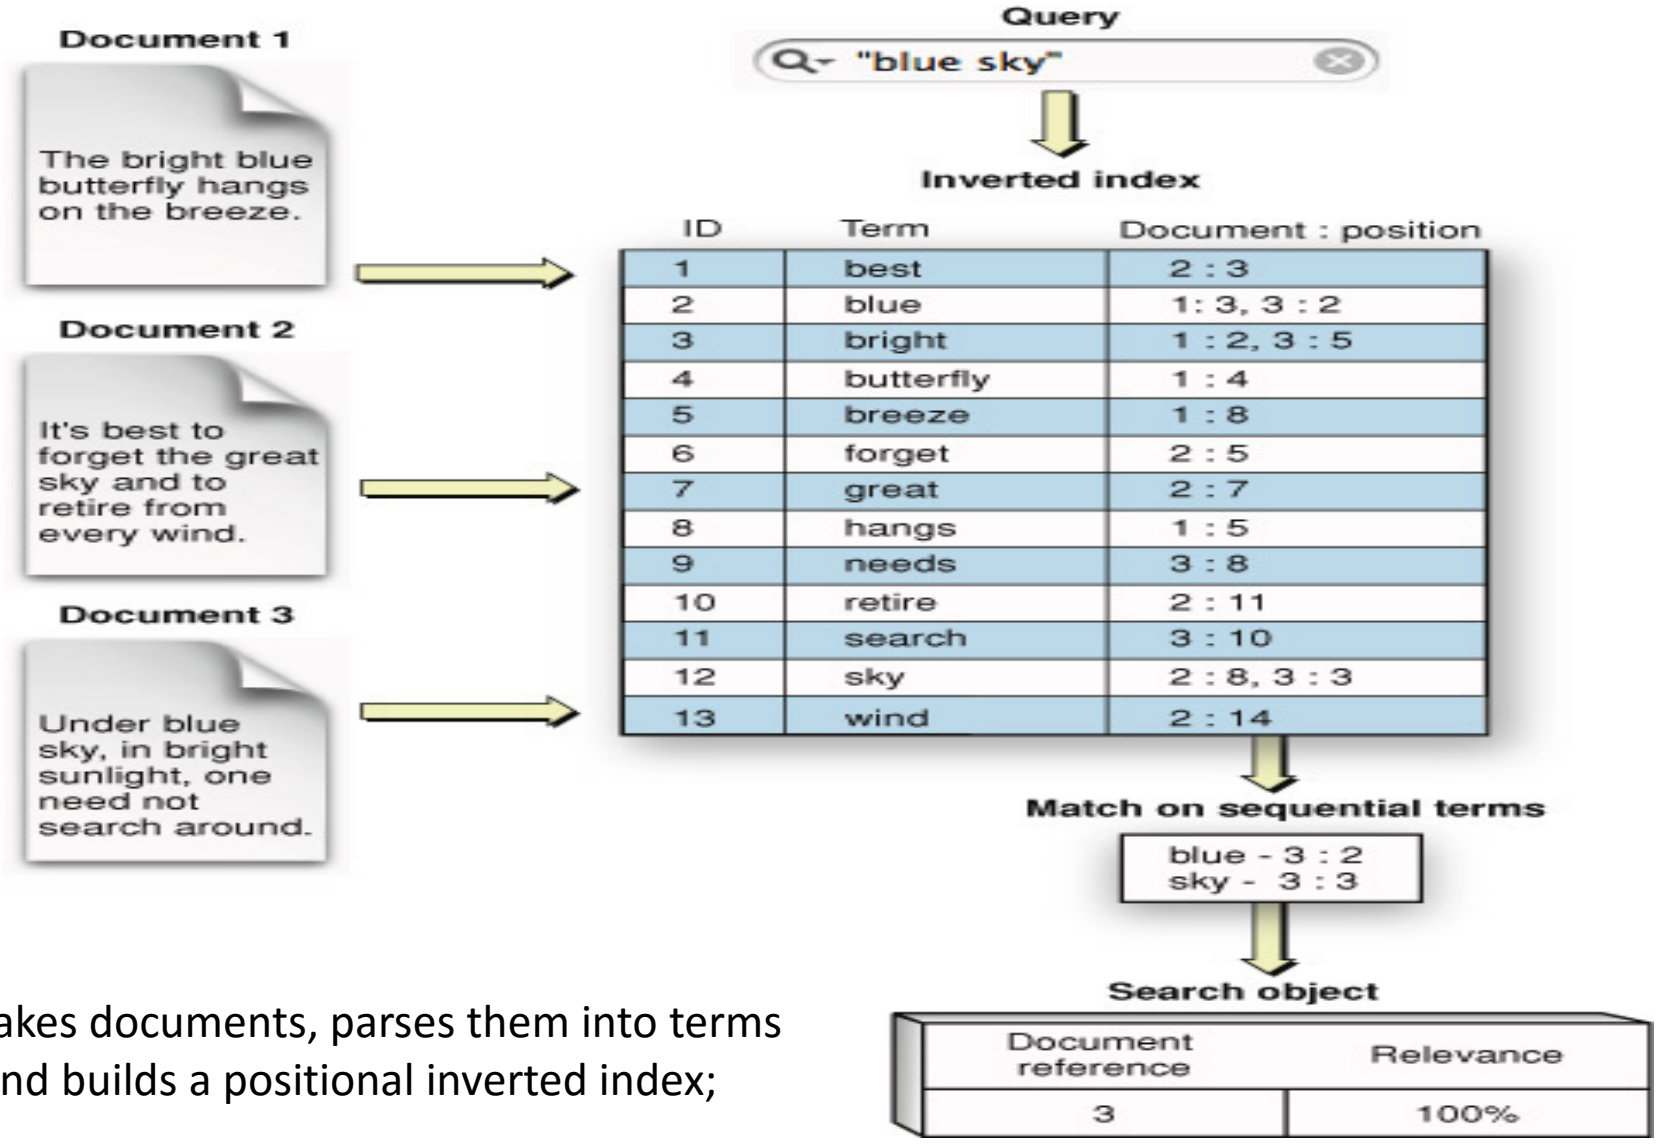

takes documents, parses them into terms and builds a positional inverted index;

Lucene Indexing Pipeline

Analyzer

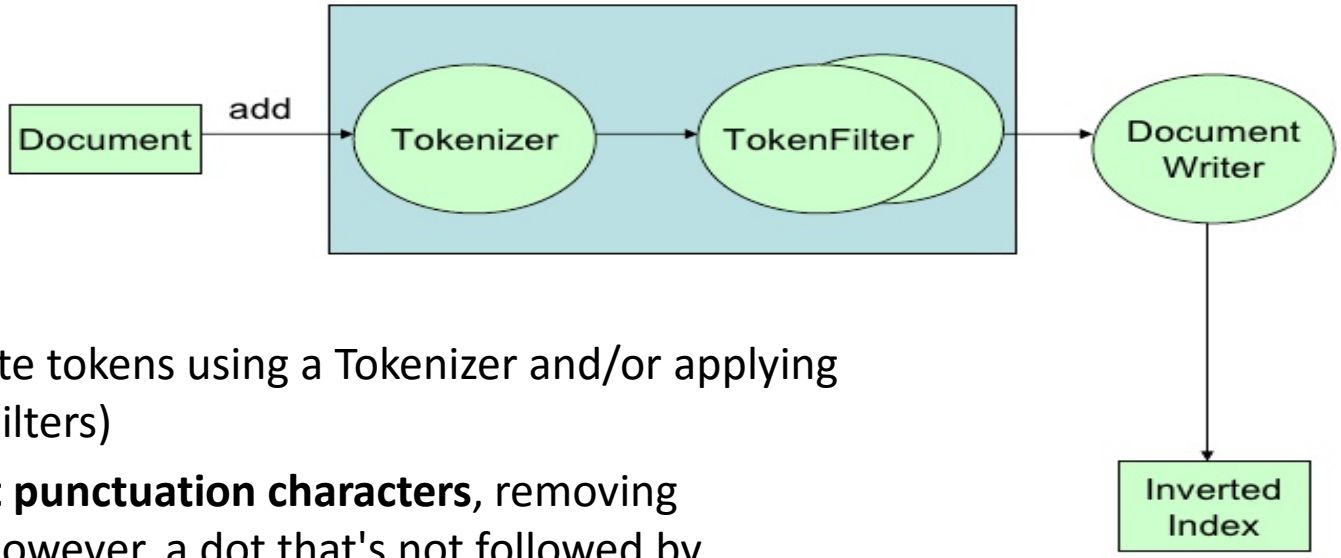

- Analyzer : create tokens using a Tokenizer and/or applying Filters (Token Filters)
- **Splits words at punctuation characters**, removing punctuation. However, a dot that's not followed by whitespace is considered part of a token.
- **Splits words at hyphens**, unless there's a number in the token, in which case the whole token is interpreted as a product number and is not split.
- **Recognizes** email addresses and internet hostnames as one token.

Analyzers, Tokenizers, and Filters

- Field analyzers are used both during ingestion, when a document is indexed, and at query time.
  - An analyzer examines the text of fields and generates a token stream.
- Tokenizers break field data into lexical units, or *tokens*.
- Filters examine a stream of tokens and keep them, transform or discard them, or create new ones.
- Analyzers: Tokenizers and filters may be combined to form pipelines, or *chains*, where the output of one is input to the next. Such a sequence of tokenizers and filters is called an *analyzer* and the resulting output of an analyzer is used to match query results or build indices.

Lucene Scoring Concepts

TF - IDF

Lucene scores using a combination of TF-IDF and vector closeness

$$w_{x,y} = tf_{x,y} \times \log\left(\frac{N}{df_x}\right)$$

TF-IDF

Term *x* within document *y*

$tf_{x,y}$ = frequency of *x* in *y*

df_x = number of documents containing *x*

N = total number of documents

TF - IDF = Term Frequency **X** Inverse Document Frequency

cosine-similarity(query_vector, document_vector) = V(q) * V(d)/|V(q)|*|V(d)|
where V(q)*V(d) is the dot product of the weighted vectors and |V(q)|, |V(d)| are the Euclidean norms of the vectors (square root of the sum of squares)

for details see
https://lucene.apache.org/core/4_0_0/core/org/apache/lucene/search/similarities/TFIDFSimilarity.html

Apache Solr

- Created by Yonik Seeley for CNET
- Enterprise Search platform for Apache Lucene
- Open source
- Highly reliable, scalable, fault tolerant
- Support distributed Indexing (SolrCloud), Replication, and load balanced querying
- **http://lucene.apache.org/solr**

High level overview

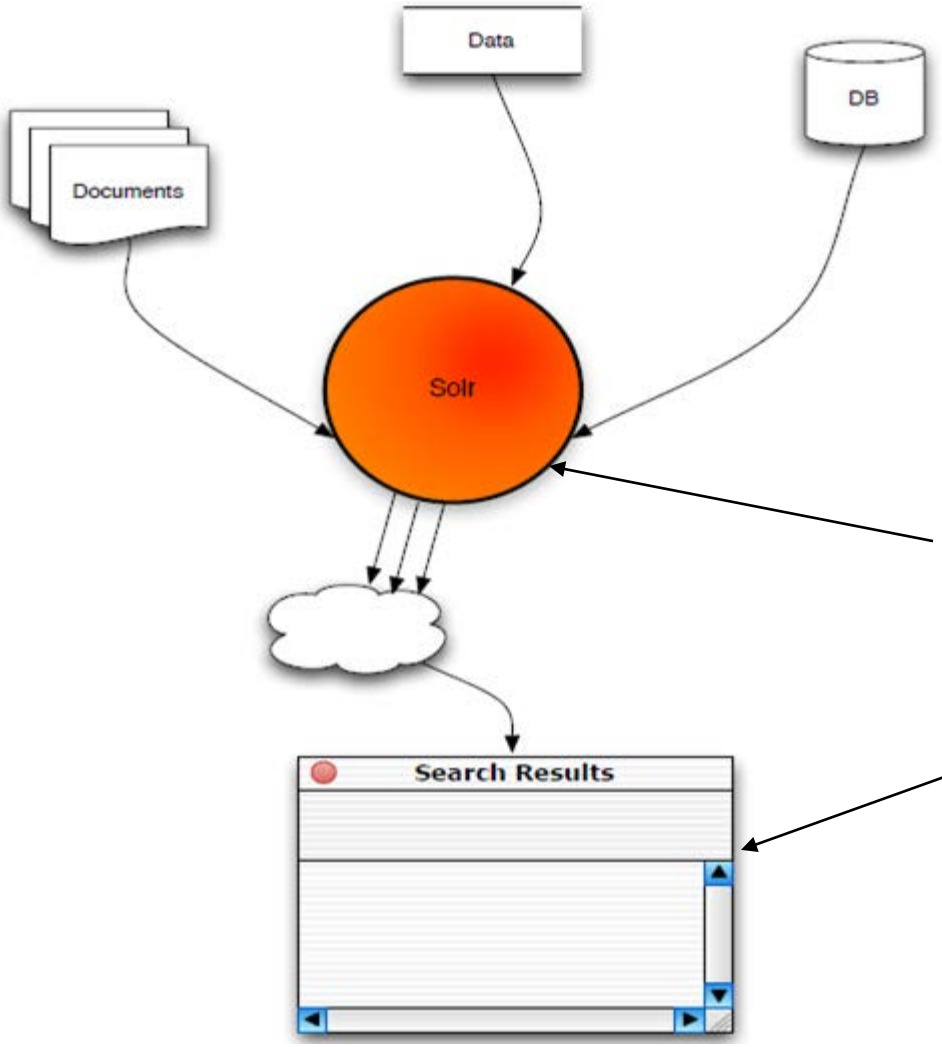

there is a Solr server which accepts queries in the form of http requests and returns results as JSON strings

Solr in Production

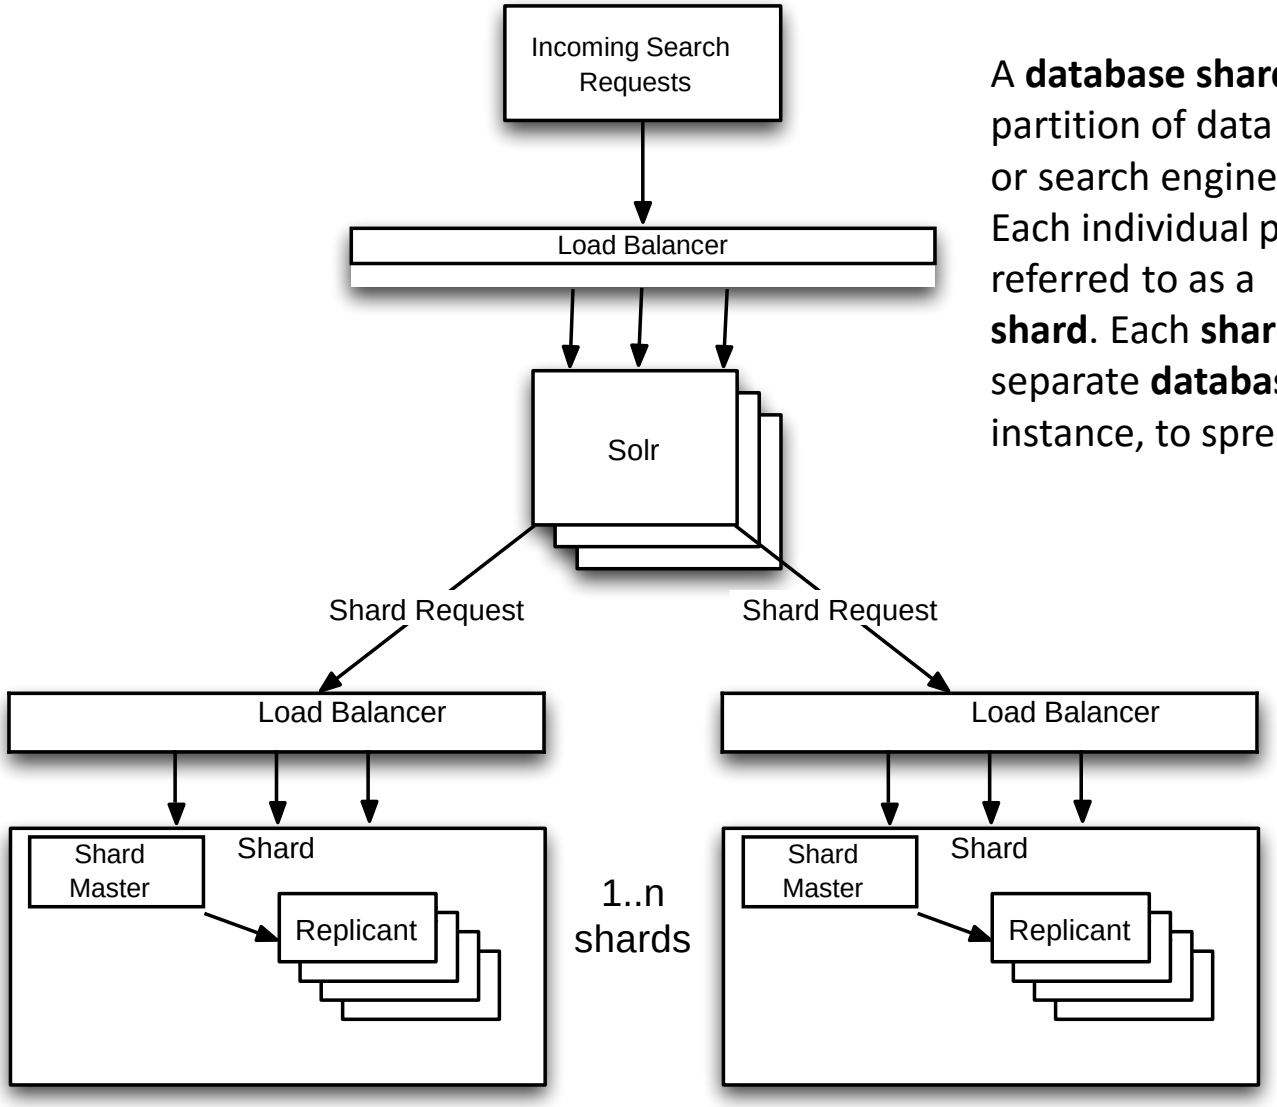

A **database shard** is a horizontal partition of data in a **database** or search engine.
Each individual partition is referred to as a **shard**. Each **shard** is held on a separate **database** server instance, to spread load.

How to start

1. Start Solr

java -jar start.jar

2. Index your data

java -jar post.jar *.xml

3. Search **http://localhost:8983/solr**

localhost indicates the Solr server is running locally on port 8983

Complete installation instructions can be found at

http://www-scf.usc.edu/~csci572/2016Spring/hw3/SolrInstallation.pdf

Querying Data

HTTP GET or POST with parameters are used to specify queries

E.g. here are 4 sample queries, some with various parameters

http://solr/select?q=electronics

http://solr/select?q=electronics&sort=price+desc

http://solr/select?q=electronics&rows=50&start=50

http://solr/select?q=electronics&fl=name+price

limit results to fields: name and price

Querying Data: Results

Canonical response format is XML, though JSON is often used as well

```
<response>
  <lst name="responseHeader">
    <int name="status">0</int>
    <int name="QTime">1</int>
  </lst>
  <result name="response" numFound="14" start="0">
    <doc>
      <arr name="cat">
        <str>electronics</str>
        <str>connector</str>
      </arr>
      <arr name="features">
         <str>car power adapter, white</str>
      </arr>
      <str name="id">F8V7067-APL-KIT</str>
      ...
```

Query Types

- Single and multi term queries
  - ex fieldname:value or title: software engineer
- +, -, AND, OR NOT operators.
  - ex. title: (software AND engineer)
- Range queries on date or numeric fields,
  - ex: timestamp: [* TO NOW] or price: [1 TO 100]
- Boost queries:
  - e.g. title:Engineer ^1.5 OR text:Engineer
- Fuzzy search : is a search for words that are similar in spelling
  - e.g. roam~0.8 => noam
- Proximity Search : with a sloppy phrase query. The close together the two terms appear, higher the score.
  - ex “apache lucene”~20 : will look for all documents where “apache” word occurs within 20 words of “lucene”

Solr is Used by Many

- **Search Engine**
  - Yandex.ru, DuckDuckGo.com
- **News Paper**
  - Guardian.co.uk
- **Music/Movies**
  - Apple.com, Netflix.com
- **Events**
  - Stubhub.com, Eventbrite.com
- **Cloud Log Management**
  - Loggly.com
- **Others**
  - Whitehouse.gov
- **Jobs**
  - Indeed.com, Simplyhired.com, Naukri.com
- **Auto**
  - AOL.com
- **Travel**
  - Cleartrip.com
- **Social Network**
  - Twitter.com, LinkedIn.com, mylife.com

Lucene/Solr Architecture

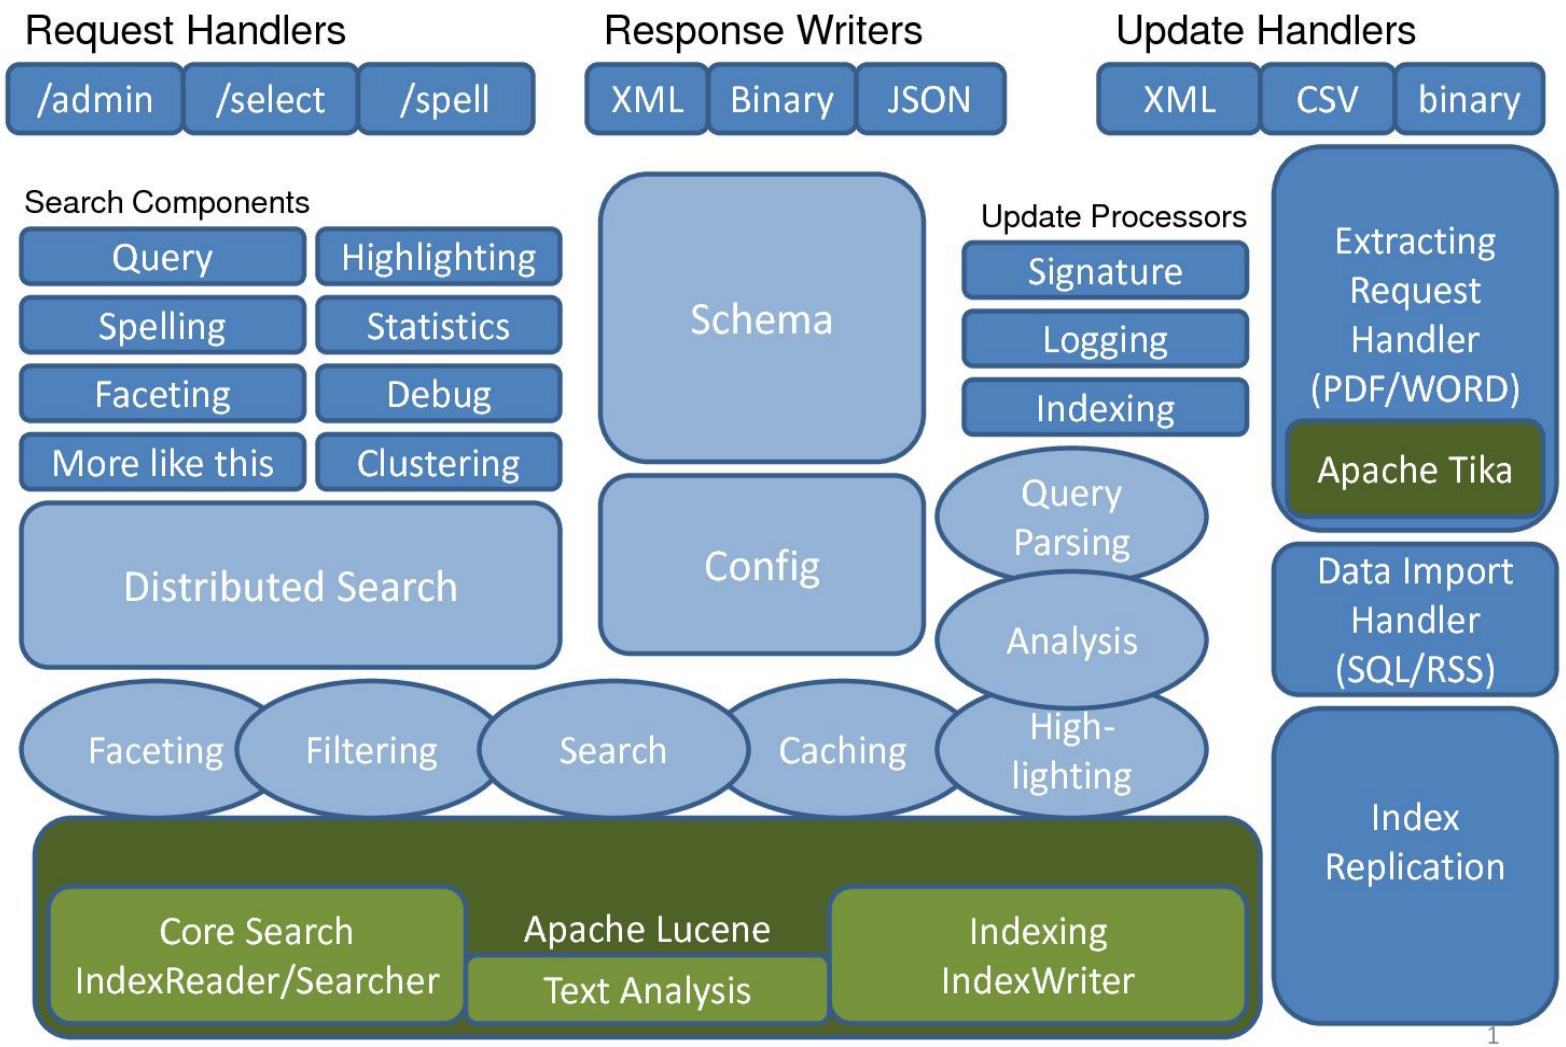

Solr Includes Spell Checking

- Not enabled by default, see example config to wire it in
- https://cwiki.apache.org/confluence/display/solr/Spell+Checking
- File or index-based dictionaries for spell correction
- Supports pluggable distance algorithms:
  - Levenstein alg: https://en.wikipedia.org/wiki/Levenshtein_distance
  - JaroWinkler alg: , https://en.wikipedia.org/wiki/Jaro%E2%80%93Winkler_distance
- http://wiki.apache.org/solr/SpellCheckComponent is a full discussion of the spell checking abilities of Solr

Solr Includes Autosuggestion

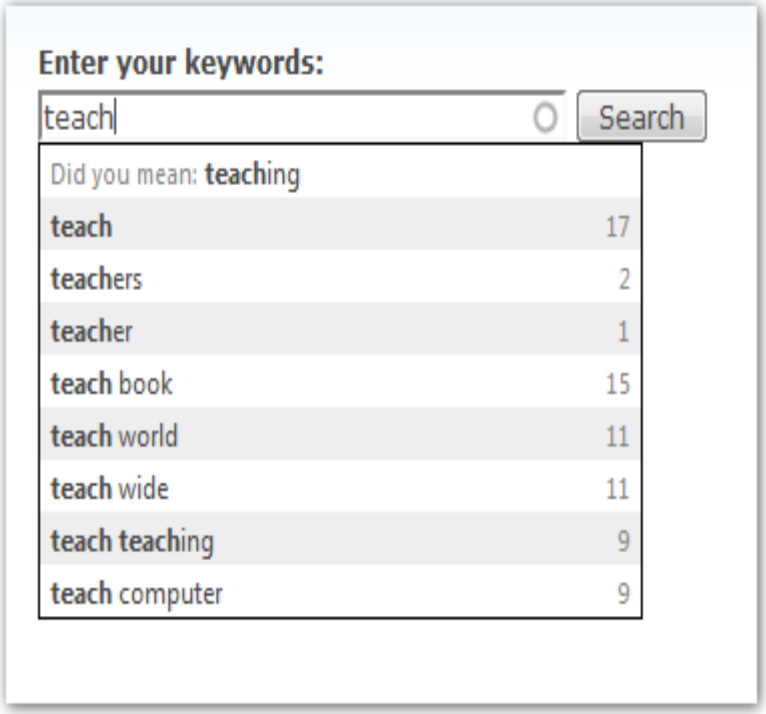

https://wiki.apache.org/solr/Suggester

Step 3: Solr to Index a Web Site

- start the Solr server
- start a new Solr core
- Use Tika to import your saved files
- Use the Solr interface to check that the files have been properly indexed
- Note the URL:

localhost:8983/solr/#/myexample

- 1413 docs successfully indexed

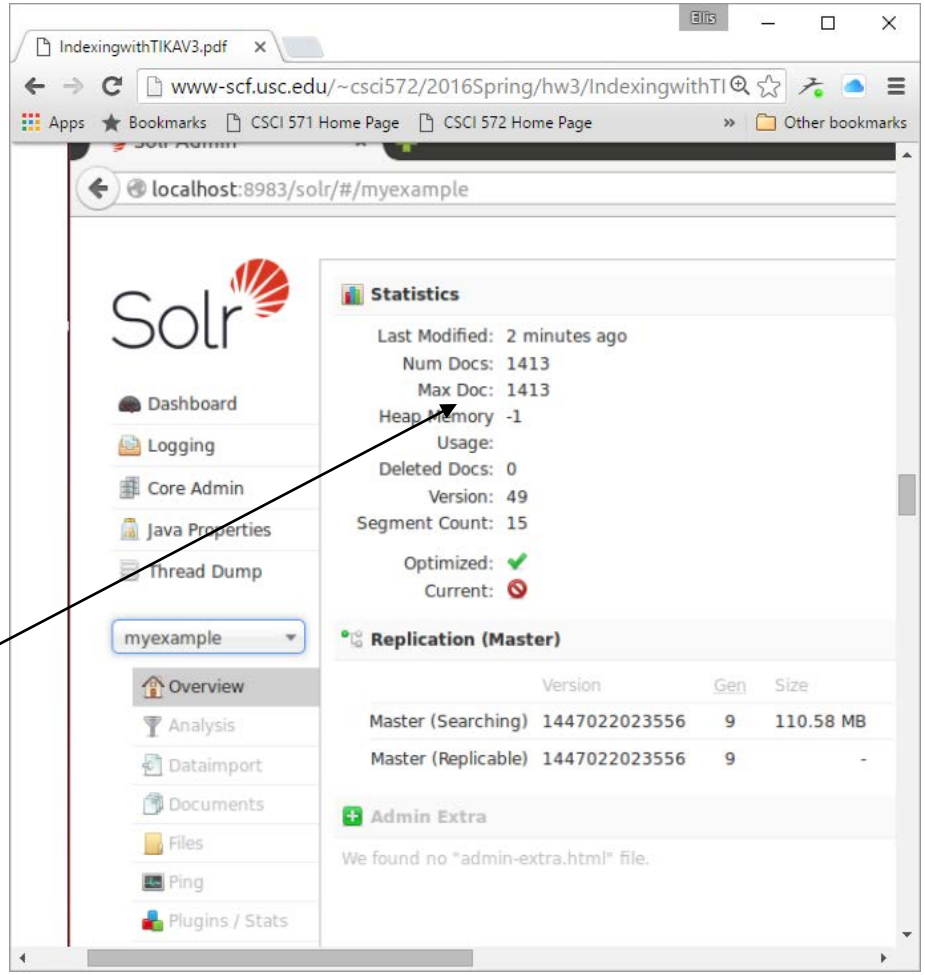

Step 4: The Actual Exercise
Comparing Search Engine Ranking Algorithms

1. From homework 2 you should have saved all of the HTML, etc. files you downloaded
2. You should install Solr as described previously
3. Take the pages you saved and index them in Solr, as described in Step 3
4. Build a front end to Solr that permits a visitor to enter a query and get matching results
5. Solr will return the results in JSON format; your server needs to take the results and format them for the user

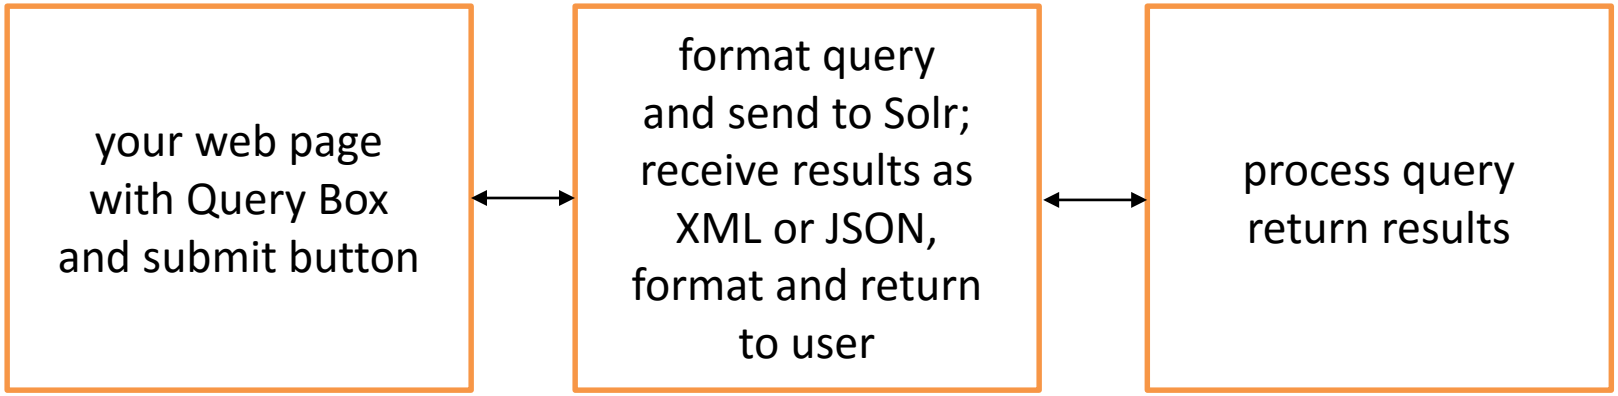

Step 4: The Actual Exercise

a PHP client that accepts input from the user in a HTML form, and sends the request to the Solr server. After the Solr server processes the query, it returns the results which is parsed by the PHP program and formatted for display

```
<?php
// make sure browsers see this page as utf-8 encoded HTML
header('Content-Type: text/html; charset=utf-8');
$limit = 10;
$query = isset($_REQUEST['q']) ? $_REQUEST['q'] : false;
$results = false;
if ($query)
{  require_once('Apache/Solr/Service.php');
 // create a new solr service instance - host, port, and corename
 // path (all defaults in this example)
 $solr = new Apache_Solr_Service('localhost', 8983, '/solr/core_name/');
 // if magic quotes is enabled then stripslashes will be needed
 if (get_magic_quotes_gpc() == 1)
 {    $query = stripslashes($query);   }
```

returning a web page

test for a query

this is the Solr client library

Solr runs on port 8983

handles quoting of special characters in query

PhP Program (2 of 3)

```
try
{    $results = $solr->search($query, 0, $limit);   }
catch (Exception $e)
{  die("<html><head><title>SEARCH EXCEPTION</title><body><pre>{$e-
>__toString()}</pre></body></html>");    }    }
?>
<html>   <head>   <title>PHP Solr Client Example</title>   </head>   <body>
 <form accept-charset="utf-8" method="get">
 <label for="q">Search:</label>
 <input id="q" name="q" type="text" value="<?php echo htmlspecialchars($query, ENT_QUOTES,
'utf-8'); ?>"/>
 <input type="submit"/>
 </form>
<?php
// display results
if ($results)
{  $total = (int) $results->response->numFound; $start = min(1, $total); $end = min($limit, $total);
?>
 <div>Results <?php echo $start; ?> - <?php echo $end;?> of <?php echo $total; ?>:</div>
 <ol>
```

send query to Solr

catch any exception

create web page output

create input text box

create submit button

end form

JSON result string

PhP Program (3 of 3)

```
<?php
 // iterate result documents
 foreach ($results->response->docs as $doc)
 {   ?>   <li>
 <table style="border: 1px solid black; text-align: left">
<?php
 // iterate document fields / values
 foreach ($doc as $field => $value)
 {    ?>
 <tr>
 <th><?php echo htmlspecialchars($field, ENT_NOQUOTES, 'utf-8'); ?></th>
 <td><?php echo htmlspecialchars($value, ENT_NOQUOTES, 'utf-8'); ?></td>
 </tr>
<?php    }
?> </table>   </li>
<?php    }    ?>
 </ol>
<?php
}   ?> </body>  </html>
```

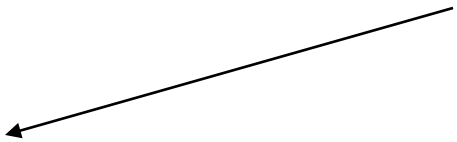

Comparing Ranking Algorithms

- we have already seen the built-in Solr ranking method
  - see slide 9
- Solr permits alternative ranking algorithms
  - we will use Page Rank as contained in an External file

http://networkx.github.io/documentation/networkx-1.10/reference/generated/networkx.algorithms.link_analysis.pagerank_alg.pagerank.html

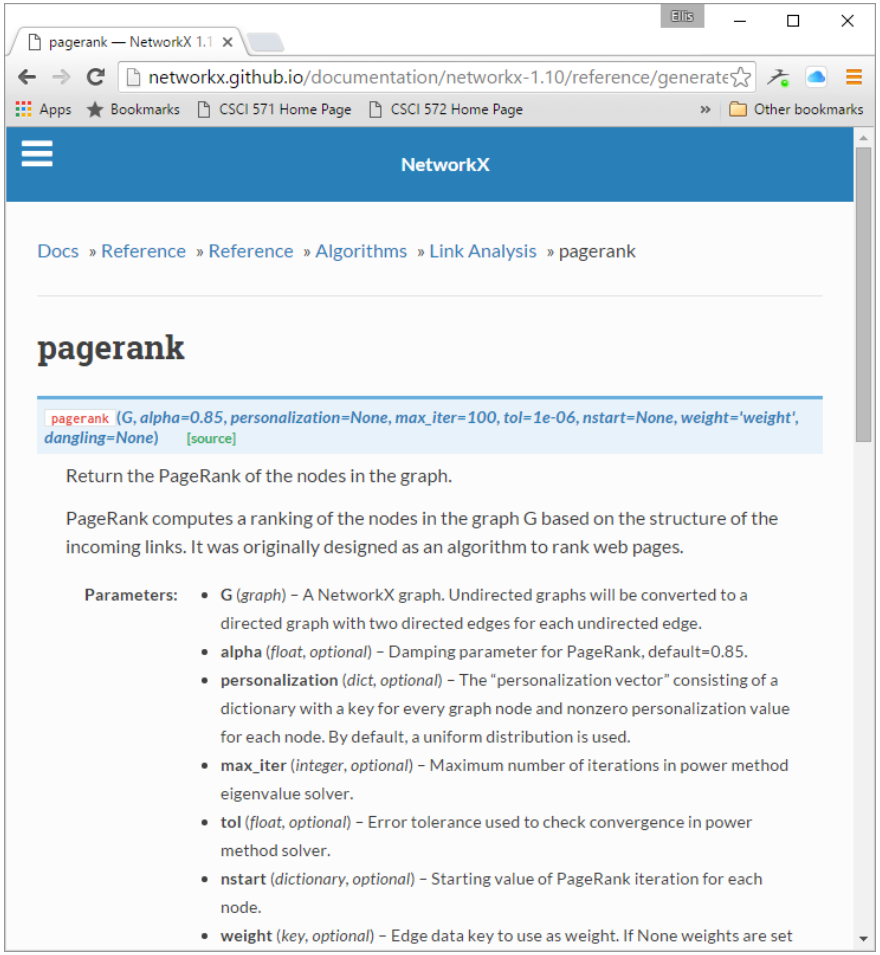

- You are going to use an open source PageRank algorithm, located at URL above;
- You should have already stored your downloaded documents with all outgoing links in a file;
- You need to create a graph that the PageRank algorithm can work on

Important Parameters:
- a NetworkX graph
- a damping parameter (e.g. 0.85)
- maximum number of iterations
- error tolerance
- starting Page Rank value of nodes

Final Steps

- Input to the PageRank algorithm is a file containing every document ID and associated with each ID, the IDs that are pointed to by links withing the document ID
- Output from the PageRank algorithm is a file containing every document ID and its associated PageRank
- place this file is solr-5.3.1/server/solr/core_name, call the file external_pageRankFile.txt
- add the PageRank field to the schema.xml file

```
<fieldType name="external" keyField="id" defVal="0" class="solr.ExternalFileField" valtype="pfloat" />
<field name="pageRankFile" type="external" stored="false" indexed="false" />
```

- Once both ranking algorithms are working you should input the same queries as Exercise #1 and compare the results